

CS 331 NP Notes 3

Graph Coloring

- Problem Definition: Given a graph $G = (V, E)$, find the smallest number of different colors to assign for each node G so that no two nodes of the same color share an edge.
 - Decision Version: Given a graph G and a bound k , does G have a k -coloring?

Why this problem is useful:


An application of this problem is when allocating resources in the presence of conflict. For example, given a set of variables and k registers. You need to map variables to registers. If two variables are use at a common point in time, they cannot be assigned to the same register. We build a graph G on the set of variables, joining two variables by an edge if they are both in use at the same time.

2-Coloring Problem: when $k = 2$, the graph coloring problem is straightforward to solve. We need to check whether a graph is bipartite.

Lemma: A graph G is colorable if and only if it is bipartite.

3-Coloring Problem: is a very difficult problem.


- Need to show the 3-coloring problem is \mathcal{NP} -Complete
 - Show $3\text{-Coloring} \in \mathcal{NP}$
 - Show $3\text{-SAT} \leq_p 3\text{-Coloring}$
- Use the simplest version of 3-SAT i.e. 1 clause
 - $F = (x_1 \vee x'_2 \vee x'_3)$
 - First, we need to find a way to turn this problem into a graph
 - For each variable x_i , create 2 nodes in G , one for v_i and one for v'_i and connect these two nodes by an edge.
 - Define 3 “special” nodes, T, F, and B, joined in a triangle (you can think of T as the color True, F as the color False, and B as the color Base).
 - Connect every variable to B. (see graph below)



- Note that this graph implicitly determines a truth assignment for the variables in the 3-SAT problem.
- The graph above has some useful properties
 - Each v_i and v'_i must get different colors

- Each v_i and v'_i must get a different color from B
 - T, F , and B must get different colors
 - if x_i is set to 1 in the 3-SAT formula, the node v_i will get the T color.
 - Hence, any 3-Coloring of the nodes of G implicitly defines a valid truth assignment for the 3-SAT instance
- We need to grow G so that only satisfying assignments can be extended to the 3-coloring of the full graph.
 - Consider the following clause: $x_1 \vee x'_2 \vee x_3$, in the language of 3-coloring of G , it says that:

"At least one of the nodes v_1, v'_2, v_3 should get the T color. "
 - We need to construct a graph with this property
 - Need to modify our graph G with a "gadget" to form G' . (See Next Page).



- Note: line color does not denote coloring of vertices; it is simply there for readability.
- The six-node "gadget" attaches to the rest of G at five existing nodes: True, False, v_1, v'_2 , and v_3 .
 - This ensures that in the event that neither x_1, x'_2 , or x_3 are assigned true, that the lowest two shaded nodes in the subgraph MUST receive the B color, the three shaded nodes above them must receive, respectively, the F, B, and T colors, and hence there's no color that can be assigned to the topmost shaded node.

- Meaning, that if no valid assignment is possible, then there will be no valid 3-Coloring for G' .

- **Proof:**

- Show $3\text{-Coloring} \in \mathcal{NP}$
 - If we are given 3 subsets and told that no edges exist between the 3 this is easy to check.
 - For every node simply check all neighbors and ensure that no neighbor exists in the same subset to which it itself is a member.
 - It takes $O(n^2)$ to verify a graph of n nodes in this fashion
- Show $3\text{-SAT} \leq_p 3\text{-Coloring}$
 - For a larger problem, start with the graph G such that there is a pair of nodes x_i, x'_i for each $x_i \in X$ connected by edges to a base node B . Then, for each clause in the 3-SAT problem attach a six-node gadget to G to form G'
 - Claim: The given 3-SAT instance is satisfiable if and only if G' has a 3-coloring.
 - Suppose there is a satisfying assignment for the 3-SAT instance. We define a coloring of G' by first coloring the B, T, and F nodes arbitrarily with the 3 given colors. Next, for each i assign v_i the color of T if $x_i = 1$ and assign v_i the color of F if $x_i = 0$. Next assign each v'_i the only remaining color such that no conflict exists. This assignment should be able to be extended into G' by the reasoning mentioned in the setup.

Conversely, suppose G' has a 3-coloring. In this coloring each node v_i is assigned either the true color or the false color; and we set the literal in F correspondingly. We now claim that in each clause of the 3-SAT instance, at least one of the terms in the clause has a truth value of one; otherwise then all 3 of the nodes in the clause would have to have the false coloring. However, as explained in the setup this would result in a graph G' which was NOT 3-Colorable which is a contradiction to our original claim. ■


Approximation Algorithms

- Given an \mathcal{NP} -Complete problem to solve, no polynomial time algorithm exists to provide an optimal solution
 - An algorithm must be:
 - Deterministic
 - Always correct
 - Bounded by a poly. time function of its input size.
- Can we develop a polynomial time algorithm that is guaranteed to provide a close to optimal solution?
- Challenge: Need to prove a solution's value is close to optimal without knowing the optimum value.

Greedy Algorithms and Bounds on the Optimum: A Load Balancing Problem


- Problem: Given M machines m_1, m_2, \dots, m_M and n jobs where each job j has a processing time t_j .
 - Need to assign each job to a machine
 - Need to balance the loads across all machines
- Formulation
 - Let $A(i)$ be the set of jobs assigned to M_i .
 - The total load on M_i : $T_i = \sum_{j \in A(i)} t_j$ where t_j = the processing time for a job
 - We want to minimize the *makespan* (the maximum load on any machine)
 $T = \max_i T_i$
- Example 1:
 - $n = \{2, 3, 4, 6, 2, 2\}$
 - $M = \{m_1, m_2, m_3\}$
 - Greedy Algorithm: Always place the job on the least loaded machine
 - However, this does NOT provide the optimal solution in this case

Greedy Solution



$$T = 8$$

Optimal Solution



$$T = 7$$

- The greedy algorithm
 - Start with no jobs assigned
 - Set $T_i = 0$ and $A(i) = \emptyset$ for all m_i
 - For $j = 1, \dots, n$
 - Let m_i be a machine that achieves the minimum $\min_k T_k$
 - Assign job j to machine m_i
 - Set $A(i) \leftarrow A(i) \cup \{j\}$
 - Set $T_i \leftarrow T_i + T_j$
 - EndFOR

Proving the Greedy Algorithm is a Correct Approximation

- Need to show the greedy solution is no more than a factor of 2 from the optimal solution
- Therefore...

- Let T = The makespan of our greedy solution
- Let T^* = The makespan of the optimal solution
 - We know that the optimal makespan must be at least:

$$T^* \geq \frac{1}{m} \sum_j t_j$$

Since the value of the makespan equals the max value over all machines.

- However, this formula is not very useful if you have one extremely long job in relation to all the shorter jobs.
 - In this case the greedy solution actually provides the optimal solution
 - Need a different lower bound for T^* to reflect this however
 - In this case the optimal makespan is at least

$$T^* \geq \max_j t_j$$
- Note: We need the two formulas mentioned above for our proof

- Claim: Our greedy algorithm produces an assignment of n jobs to M machines within a factor of 2 of the optimal solution; i.e. $T \leq 2T^*$
 - **Proof**:
 - Consider M_i attains the max load T in the assignment
 - Consider the last job j that was placed on M_i .




Figure 11.2 Accounting for the load on machine M_i in two parts: the last job to be added, and all the others.

- When we assign job j to M_i , M_i had the smallest load of all machines $T_i - t_j$ before adding j .
- Every machine had a load at least $T_i - t_j$
- Add up the load on all machines and we obtain
$$\sum_k t_k \geq m(T_i - t_j) \equiv \frac{1}{m} \sum_k t_k \geq T_i - t_j$$
- Now apply the two formulas above to form a sequence of equations

$T_i - t_j \leq \frac{1}{m} \sum_k t_k \leq T^*$ (when t_j is **not an** extremely long job)

$$T_i - t_j \leq T^* \quad (1)$$

$$t_j \leq T^* \quad (2) \text{ (when } t_j \text{ is an extremely long job)}$$

By adding (1) and (2), we get the following.

$$T_i \leq 2T^*$$

Therefore,

$T_i \leq 2T^* \Rightarrow T \leq 2T^*$ by our assumption for the proof that T_i is the machine with the maximum makespan ■

But wait,
there's more!



- We can optimize the algorithm by sorting jobs in decreasing order first
 - Optimized greedy algorithm
 - Start with no jobs assigned
 - Set $T_i = 0$ and $A(i) = \emptyset$ for all machines m_i
 - Sort jobs in decreasing order of processing time t_j
 - Assume that $t_1 \geq t_2 \geq \dots \geq t_n$
 - For $j = 1, \dots, n$
 - Let m_i be a machine that achieves the minimum $\min_k T_k$
 - Assign job j to machine m_i
 - Set $A(i) \leftarrow A(i) \cup \{j\}$
 - Set $T_i \leftarrow T_i + t_j$
 - EndFOR
- However, just saying this is better is not enough, we need to prove it
 - More specifically, we need to prove that $T \leq \frac{3}{2}T^*$
- First however, we need to prove that doing this will produce an optimal solution
 - Lemma: If there are more than M jobs, then $T^* \geq 2t_{M+1}$
 - Consider the first $M + 1$ jobs in decreasing order
 - Each job MUST need at least t_{M+1} time to process
 - There are $M + 1$ jobs but only M machines
 - \therefore it must be the case that one machine has at least two jobs
 - \therefore it must be the case that the processing time of the machine will be at least $2t_{m+1}$
- Proving $T \leq \frac{3}{2}T^*$
 - Consider M_i attains the max load T in the assignment
 - If M_i only holds a single job, then the schedule is optimal. Otherwise, assume machine M_i has at least two jobs, and let t_j be the last job assigned to the machine
 - It must be the case that $j \geq M + 1$ since the algorithm assigns the first M jobs to M distinct machines
 - Thus $t_j \leq t_{m+1}$ (jobs are sorted in decreasing order)
 - Therefore, $t_j \leq t_{m+1} \leq \frac{1}{2}T^* \leftarrow$ occurs because of what we proved above ($T^* \geq 2t_{m+1} \equiv \frac{1}{2}T^* \geq t_{m+1}$)
 - Therefore, repeating the inequalities from the first proof we now have

$$T_i - t_j \leq \frac{1}{m} \sum_k t_k \leq T^*$$

$$T_i - t_j \leq T^* \quad (1)$$

$$t_j \leq \frac{1}{2}T^* \quad (2)$$

Adding (1) and (2), we get the following.

$$T_i \leq \frac{3}{2}T^* \Rightarrow T \leq \frac{3}{2}T^* \blacksquare$$